

SPOTIFY DATA ANALYTICS

Objective:

Analyze Spotify data to uncover insights about music trends, user preferences, and song characteristics.

```
# importing major libraries
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
import plotly.express as px
from textblob import TextBlob

# addition libraries
import warnings
warnings.filterwarnings('ignore')
```

Task 1.1: Overview of the Project and Its Objectives

This project focuses on analyzing Spotify music data to uncover meaningful patterns, trends, and insights related to songs, artists, genres, and listener preferences. With the rapid growth of digital music streaming, platforms like Spotify generate vast amounts of data that can be used to understand what makes songs popular, how musical styles evolve over time, and how user tastes vary across genres and artists.

The main goal of this project is to apply exploratory data analysis (EDA) and data-driven techniques to examine the characteristics of tracks and identify relationships between audio features and popularity. By analyzing attributes such as danceability, energy, tempo, loudness, and valence, the project aims to understand how musical properties influence audience engagement and trends.

Objectives

- Explore and clean the Spotify datasets for analysis
- Analyze track-level audio features and their distributions
- Identify trends across artists, genres, and years
- Examine how musical characteristics affect popularity
- Discover patterns in listener preferences and genre performance
- Generate insights and recommendations using data visualization and statistical analysis

Overall, the project demonstrates how data analytics can be used to better understand the music industry and listener behavior.

Task 1.2: Explanation of the Dataset Used

The project uses multiple Spotify datasets containing detailed information about tracks, artists, genres, and yearly trends. These datasets provide both musical attributes and metadata that support comprehensive analysis.

Datasets Included

1. Track-level Dataset (data.csv)

This dataset contains information about individual songs available on Spotify. Each row represents a track with various audio and popularity features.

Key attributes include:

- Track name and artist
- Popularity score
- Duration
- Danceability
- Energy
- Loudness
- Speechiness
- Acousticness
- Instrumentalness
- Liveness
- Valence (musical positivity)
- Tempo

This dataset is primarily used to analyze how audio characteristics influence song popularity and listener engagement.

2. Artist-level Dataset (data_by_artist.csv)

This dataset aggregates features at the artist level by averaging track characteristics for each artist.

Used to:

- Compare artists based on their musical style
- Identify high-performing or popular artists
- Analyze consistency in audio features across an artist's songs

3. Genre-level Dataset (data_by_genres.csv)

This dataset groups tracks by genre and provides averaged audio features.

Used to:

- Study differences between musical genres
- Identify dominant genre traits (e.g., energetic vs acoustic genres)
- Compare genre popularity

4. Yearly Dataset (data_by_year.csv)

This dataset summarizes music trends over time by year.

Used to:

- Analyze how music characteristics evolve across years
- Observe changes in popularity, tempo, or energy over time
- Detect historical trends in the music industry

5. Tracks with Genres Dataset (data_w_genres.csv)

This dataset links tracks with their associated genres, enabling more detailed analysis.

Used to:

- Connect individual tracks to genre categories
- Perform combined track + genre analysis
- Improve genre-based recommendations and comparisons

Summary

Together, these datasets provide a comprehensive view of Spotify music data at multiple levels — track, artist, genre, and time. This multi-dimensional structure enables deeper insights into musical patterns, listener preferences, and evolving trends in the music industry.

```
tracks = pd.read_csv('data.csv')
artist = pd.read_csv('data_by_artist.csv')
genres = pd.read_csv('data_by_genres.csv')
year = pd.read_csv('data_by_year.csv')
tracks_genres = pd.read_csv('data_w_genres.csv')
```

Data Collection and Preprocessing

Task 2.1: Load the dataset into the Jupyter notebook.

Task 2.2: Inspect and clean the data (handle missing values, duplicates, etc.).

Task 2.3: Perform exploratory data analysis (EDA) to understand the basic characteristics of the data

```
artist.head()
```

```
      mode  count  acousticness \
0       1      9     0.590111
1       1     26     0.862538
2       1      7     0.856571
3       1     27     0.884926
4       1      7     0.510714

                                                artists  danceability \
0          "Cats" 1981 Original London Cast     0.467222
1          "Cats" 1983 Broadway Cast     0.441731
2    "Fiddler On The Roof" Motion Picture Chorus     0.348286
3  "Fiddler On The Roof" Motion Picture Orchestra     0.425074
4 "Joseph And The Amazing Technicolor Dreamcoat"...

      duration_ms   energy  instrumentalness  liveness  loudness \
0  250318.555556  0.394003           0.011400  0.290833 -14.448000
1  287280.000000  0.406808           0.081158  0.315215 -10.690000
2  328920.000000  0.286571           0.024593  0.325786 -15.230714
3  262890.962963  0.245770           0.073587  0.275481 -15.639370
4  270436.142857  0.488286           0.009400  0.195000 -10.236714

  speechiness        tempo  valence  popularity  key
0  0.210389  117.518111  0.389500  38.333333  5
1  0.176212  103.044154  0.268865  30.576923  5
2  0.118514   77.375857  0.354857  34.857143  0
3  0.123200   88.667630  0.372030  34.851852  0
4  0.098543  122.835857  0.482286  43.000000  5
```

```
genres.head()
```

```
      mode          genres  acousticness  danceability
duration_ms \
0       1  21st century classical     0.979333     0.162883
1.602977e+05
1       1            432hz     0.494780     0.299333
1.048887e+06
2       1            8-bit     0.762000     0.712000
1.151770e+05
3       1              []     0.651417     0.529093
2.328809e+05
4       1        a cappella     0.676557     0.538961
1.906285e+05

      energy  instrumentalness  liveness  loudness  speechiness
tempo \
0  0.071317           0.606834  0.361600 -31.514333     0.040567
75.336500
1  0.450678           0.477762  0.131000 -16.854000     0.076817
```

```

120.285667
2 0.818000          0.876000  0.126000 -9.180000      0.047000
133.444000
3 0.419146          0.205309  0.218696 -12.288965      0.107872
112.857352
4 0.316434          0.003003  0.172254 -12.479387      0.082851
112.110362

    valence  popularity  key
0  0.103783   27.833333    6
1  0.221750   52.500000    5
2  0.975000   48.000000    7
3  0.513604   20.859882    7
4  0.448249   45.820071    7

year.head()

    mode  year  acousticness  danceability  duration_ms  energy \
0     1  1921       0.886896      0.418597  260537.166667  0.231815
1     1  1922       0.938592      0.482042  165469.746479  0.237815
2     1  1923       0.957247      0.577341  177942.362162  0.262406
3     1  1924       0.940200      0.549894  191046.707627  0.344347
4     1  1925       0.962607      0.573863  184986.924460  0.278594

    instrumentalness  liveness  loudness  speechiness  tempo
valence \
0            0.344878  0.205710 -17.048667      0.073662  101.531493
0.379327
1            0.434195  0.240720 -19.275282      0.116655  100.884521
0.535549
2            0.371733  0.227462 -14.129211      0.093949  114.010730
0.625492
3            0.581701  0.235219 -14.231343      0.092089  120.689572
0.663725
4            0.418297  0.237668 -14.146414      0.111918  115.521921
0.621929

    popularity  key
0     0.653333    2
1     0.140845   10
2     5.389189    0
3     0.661017   10
4     2.604317    5

tracks_genres.head()

    genres
artists \
0  ['show tunes']           "Cats" 1981 Original London Cast
1                  []           "Cats" 1983 Broadway Cast

```

```

2          []      "Fiddler On The Roof" Motion Picture Chorus
3          []      "Fiddler On The Roof" Motion Picture Orchestra
4          []  "Joseph And The Amazing Technicolor Dreamcoat"...

    acousticness  danceability      duration_ms      energy
instrumentalness \
0      0.590111      0.467222  250318.555556  0.394003
0.011400
1      0.862538      0.441731  287280.000000  0.406808
0.081158
2      0.856571      0.348286  328920.000000  0.286571
0.024593
3      0.884926      0.425074  262890.962963  0.245770
0.073587
4      0.510714      0.467143  270436.142857  0.488286
0.009400

    liveness   loudness   speechiness      tempo      valence  popularity
key \
0  0.290833 -14.448000      0.210389  117.518111  0.389500  38.333333
5
1  0.315215 -10.690000      0.176212  103.044154  0.268865  30.576923
5
2  0.325786 -15.230714      0.118514   77.375857  0.354857  34.857143
0
3  0.275481 -15.639370      0.123200   88.667630  0.372030  34.851852
0
4  0.195000 -10.236714      0.098543  122.835857  0.482286  43.000000
5

    mode   count
0      1      9
1      1     26
2      1      7
3      1     27
4      1      7

tracks.head()

    valence   year  acousticness \
0  0.0594  1921      0.982
1  0.9630  1921      0.732
2  0.0394  1921      0.961
3  0.1650  1921      0.967
4  0.2530  1921      0.957

```

```

artists      danceability \
0  ['Sergei Rachmaninoff', 'James Levine', 'Berli...          0.279
1                                ['Dennis Day']           0.819
2  ['KHP Kridhamardawa Karaton Ngayogyakarta Hadi...  0.328
3                                ['Frank Parker']        0.275
4                                ['Phil Regan']         0.418

duration_ms   energy   explicit      id
instrumentalness \
0     831667   0.211       0  4BJqT0PrAfrxzM0xytFOIz
0.878000
1     180533   0.341       0  7xPhfUan2yNtyFG0cUWkt8
0.000000
2     500062   0.166       0  1o6I8BglA6ylDMrIELygv1
0.913000
3     210000   0.309       0  3ftBPsC5vPBKxYSee08FDH
0.000028
4     166693   0.193       0  4d6HGyGT8e121BsdKmw9v6
0.000002

key    liveness  loudness  mode \
0     10      0.665   -20.096    1
1      7      0.160   -12.441    1
2      3      0.101   -14.850    1
3      5      0.381   -9.316     1
4      3      0.229   -10.096    1

name      popularity
release_date \
0  Piano Concerto No. 3 in D Minor, Op. 30: III. ...          4
1921
1                                Clancy Lowered the Boom           5
1921
2                                Gati Bali                     5
1921
3                                Danny Boy                     3
1921
4                                When Irish Eyes Are Smiling      2
1921

speechiness      tempo
0      0.0366    80.954
1      0.4150    60.936
2      0.0339    110.339
3      0.0354    100.109
4      0.0380    101.665

print(tracks.info())
print('*'*50)
print(artist.info())

```

```

print('*'*50)
print(genres.info())
print('*'*50)
print(year.info())
print('*'*50)
print(tracks_genres.info())

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 170653 entries, 0 to 170652
Data columns (total 19 columns):
 #   Column           Non-Null Count   Dtype  
--- 
 0   valence          170653 non-null    float64
 1   year              170653 non-null    int64  
 2   acousticness      170653 non-null    float64
 3   artists            170653 non-null    object  
 4   danceability       170653 non-null    float64
 5   duration_ms        170653 non-null    int64  
 6   energy             170653 non-null    float64
 7   explicit            170653 non-null    int64  
 8   id                 170653 non-null    object  
 9   instrumentalness   170653 non-null    float64
 10  key                170653 non-null    int64  
 11  liveness            170653 non-null    float64
 12  loudness            170653 non-null    float64
 13  mode               170653 non-null    int64  
 14  name                170653 non-null    object  
 15  popularity           170653 non-null    int64  
 16  release_date         170653 non-null    object  
 17  speechiness          170653 non-null    float64
 18  tempo               170653 non-null    float64
dtypes: float64(9), int64(6), object(4)
memory usage: 24.7+ MB
None
-----
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 28680 entries, 0 to 28679
Data columns (total 15 columns):
 #   Column           Non-Null Count   Dtype  
--- 
 0   mode              28680 non-null    int64  
 1   count             28680 non-null    int64  
 2   acousticness      28680 non-null    float64
 3   artists            28680 non-null    object  
 4   danceability       28680 non-null    float64
 5   duration_ms        28680 non-null    float64
 6   energy             28680 non-null    float64
 7   instrumentalness   28680 non-null    float64
 8   liveness            28680 non-null    float64
 9   loudness            28680 non-null    float64

```

```

10 speechiness      28680 non-null  float64
11 tempo           28680 non-null  float64
12 valence          28680 non-null  float64
13 popularity       28680 non-null  float64
14 key              28680 non-null  int64
dtypes: float64(11), int64(3), object(1)
memory usage: 3.3+ MB
None
-----
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2973 entries, 0 to 2972
Data columns (total 14 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   mode              2973 non-null    int64  
 1   genres             2973 non-null    object  
 2   acousticness       2973 non-null    float64 
 3   danceability       2973 non-null    float64 
 4   duration_ms        2973 non-null    float64 
 5   energy             2973 non-null    float64 
 6   instrumentalness  2973 non-null    float64 
 7   liveness           2973 non-null    float64 
 8   loudness           2973 non-null    float64 
 9   speechiness        2973 non-null    float64 
 10  tempo              2973 non-null    float64 
 11  valence            2973 non-null    float64 
 12  popularity          2973 non-null    float64 
 13  key                2973 non-null    int64  
dtypes: float64(11), int64(2), object(1)
memory usage: 325.3+ KB
None
-----
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 14 columns):
 #   Column            Non-Null Count  Dtype  
--- 
 0   mode              100 non-null    int64  
 1   year              100 non-null    int64  
 2   acousticness       100 non-null    float64 
 3   danceability       100 non-null    float64 
 4   duration_ms        100 non-null    float64 
 5   energy             100 non-null    float64 
 6   instrumentalness  100 non-null    float64 
 7   liveness           100 non-null    float64 
 8   loudness           100 non-null    float64 
 9   speechiness        100 non-null    float64 
 10  tempo              100 non-null    float64 
 11  valence            100 non-null    float64

```

```
12 popularity          100 non-null    float64
13 key                 100 non-null    int64
dtypes: float64(11), int64(3)
memory usage: 11.1 KB
None
-----
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 28680 entries, 0 to 28679
Data columns (total 16 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   genres            28680 non-null   object  
 1   artists           28680 non-null   object  
 2   acousticness      28680 non-null   float64 
 3   danceability      28680 non-null   float64 
 4   duration_ms       28680 non-null   float64 
 5   energy            28680 non-null   float64 
 6   instrumentalness 28680 non-null   float64 
 7   liveness          28680 non-null   float64 
 8   loudness          28680 non-null   float64 
 9   speechiness       28680 non-null   float64 
 10  tempo             28680 non-null   float64 
 11  valence           28680 non-null   float64 
 12  popularity        28680 non-null   float64 
 13  key               28680 non-null   int64  
 14  mode              28680 non-null   int64  
 15  count             28680 non-null   int64  
dtypes: float64(11), int64(3), object(2)
memory usage: 3.5+ MB
None

print(tracks.isnull().sum())
print('*'*50)
print(artist.isnull().sum())
print('*'*50)
print(genres.isnull().sum())
print('*'*50)
print(year.isnull().sum())
print('*'*50)
print(tracks_genres.isnull().sum())

valence      0
year         0
acousticness 0
artists      0
danceability 0
duration_ms  0
energy       0
explicit     0
id           0
```

```
instrumentalness    0  
key                 0  
liveness            0  
loudness            0  
mode                0  
name                0  
popularity          0  
release_date        0  
speechiness         0  
tempo               0  
dtype: int64
```

```
mode                0  
count               0  
acousticness        0  
artists             0  
danceability        0  
duration_ms         0  
energy              0  
instrumentalness   0  
liveness            0  
loudness            0  
speechiness         0  
tempo               0  
valence             0  
popularity          0  
key                 0  
dtype: int64
```

```
mode                0  
genres              0  
acousticness        0  
danceability        0  
duration_ms         0  
energy              0  
instrumentalness   0  
liveness            0  
loudness            0  
speechiness         0  
tempo               0  
valence             0  
popularity          0  
key                 0  
dtype: int64
```

```
mode                0  
year                0  
acousticness        0  
danceability        0
```

```

duration_ms      0
energy           0
instrumentalness 0
liveness          0
loudness          0
speechiness       0
tempo             0
valence           0
popularity         0
key               0
dtype: int64
-----
genres            0
artists           0
acousticness      0
danceability       0
duration_ms       0
energy            0
instrumentalness 0
liveness          0
loudness          0
speechiness        0
tempo             0
valence           0
popularity         0
key               0
mode               0
count              0
dtype: int64

# Fill numeric missing values with mean
tracks.fillna(tracks.mean(numeric_only=True), inplace=True)
artist.fillna(artist.mean(numeric_only=True), inplace=True)
genres.fillna(genres.mean(numeric_only=True), inplace=True)
year.fillna(year.mean(numeric_only=True), inplace=True)

# Remove duplicates
tracks.drop_duplicates(inplace=True)
artist.drop_duplicates(inplace=True)
genres.drop_duplicates(inplace=True)
year.drop_duplicates(inplace=True)
tracks_genres.drop_duplicates(inplace=True)

# Data type corrections
tracks['year'] = tracks['year'].astype(int)
tracks['popularity'] = tracks['popularity'].astype(int)
year['year'] = year['year'].astype(int)

# Basic dataset overview
print(tracks.shape)

```

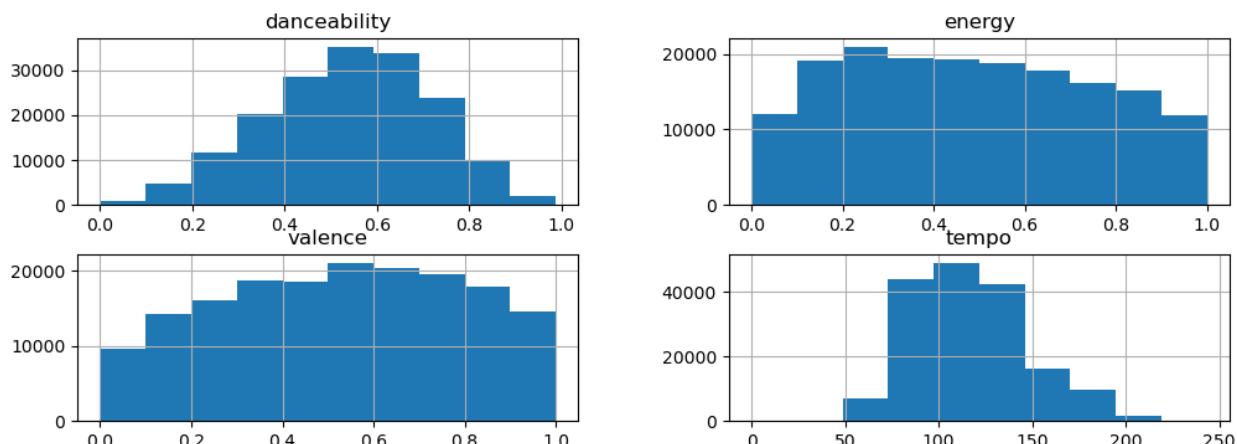
```

print('-'*50)
print(artists.shape)
print('-'*50)
print(genres.shape)
print('*'*50)
print(years.shape)
print('*'*50)
print(tracks_genres.shape)

(170653, 19)
-----
(28680, 15)
-----
(2973, 14)
-----
(100, 14)
-----
(28680, 16)

# Distribution of audio features
tracks[['danceability','energy','valence','tempo']].hist(figsize=(12,4))
plt.show()

```



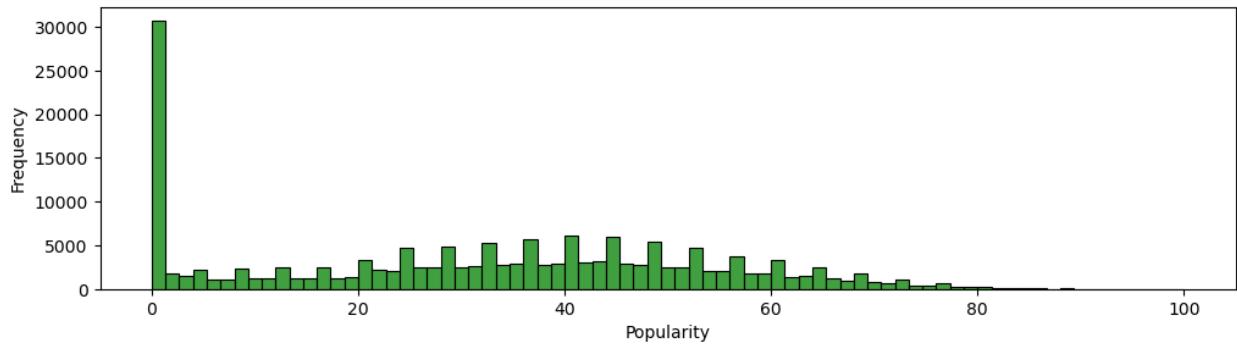
```

# Popularity analysis
print(tracks['popularity'].describe())
plt.figure(figsize=(12,3))
sns.histplot(data=tracks,x='popularity',color='green')
plt.xlabel("Popularity")
plt.ylabel("Frequency")
plt.show()

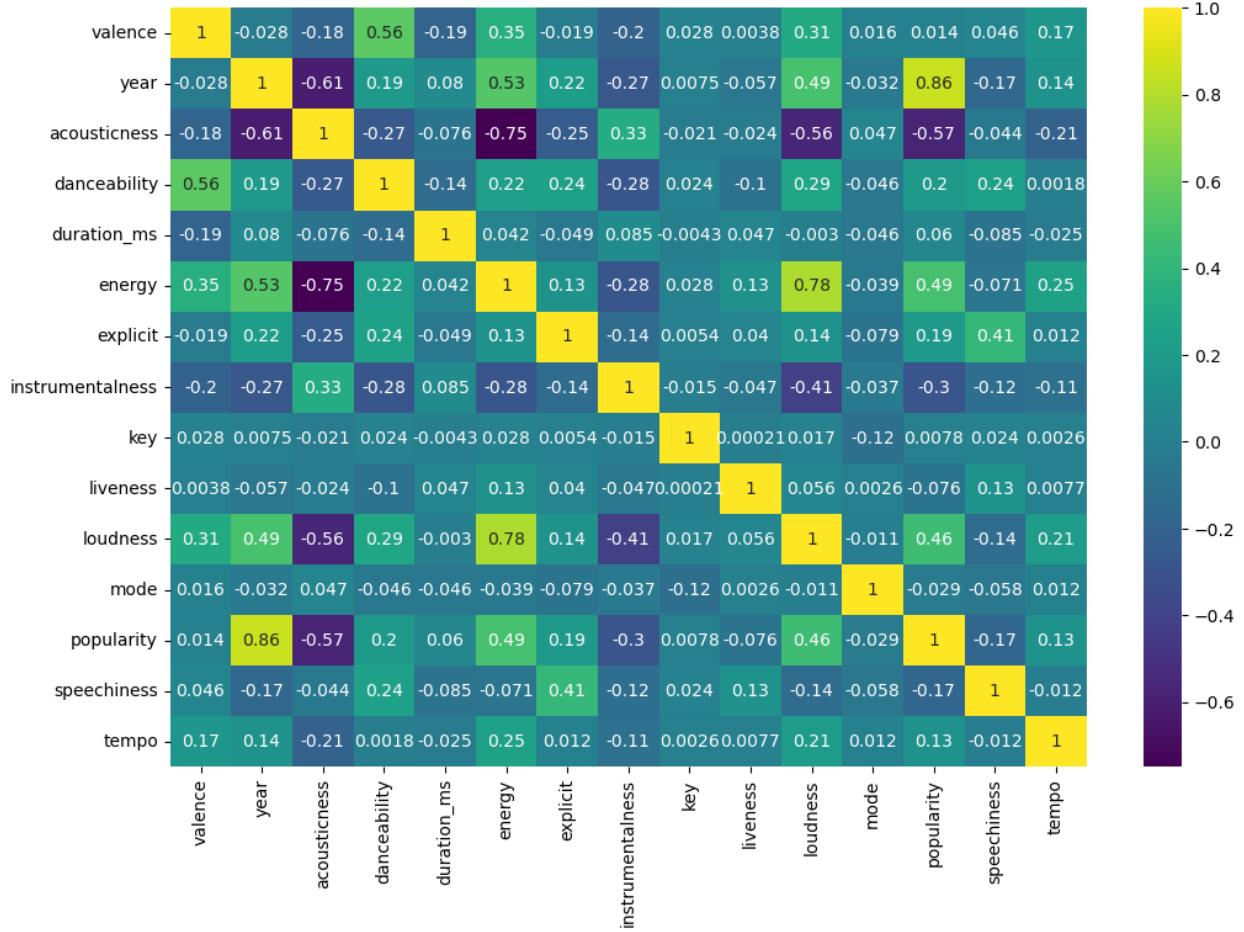
count      170653.000000
mean        31.431794
std         21.826615

```

```
min           0.000000
25%        11.000000
50%        33.000000
75%        48.000000
max       100.000000
Name: popularity, dtype: float64
```



```
# Correlation between features
plt.figure(figsize=(12,8))
sns.heatmap(tracks.corr(numeric_only=True), annot=True, cmap='viridis')
plt.show()
```



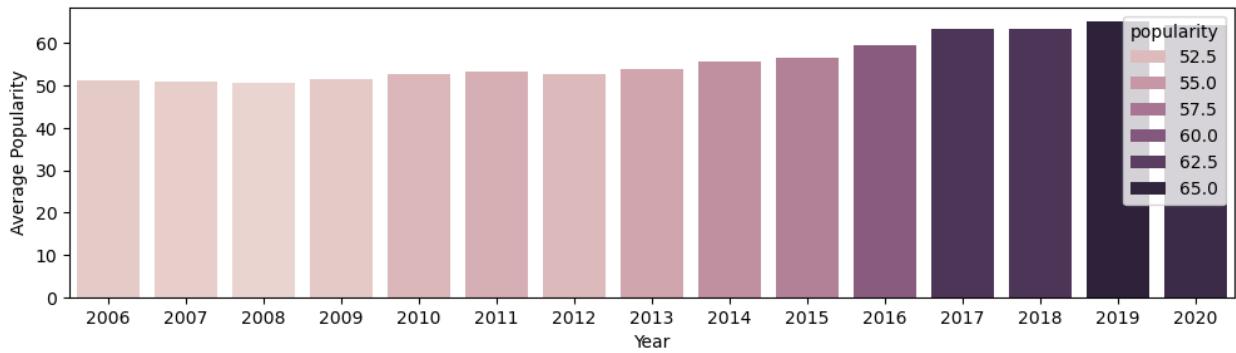
```
# Genre analysis
genres.sort_values('popularity', ascending=False).head(10)
```

	mode	genres	acousticness	danceability
duration_ms \ 237	1	basshall	0.213167	0.81800
169799.166667				
2755	1	trap venezolano	0.044600	0.87700
231848.000000				
2533	1	south african house	0.043833	0.84700
311854.333333				
2778	0	turkish edm	0.008290	0.69800
186700.000000				
46	0	alberta hip hop	0.330000	0.88500
144000.000000				
536	0	chinese electropop	0.002570	0.66000
217088.000000				
37	0	afroswing	0.318450	0.71175
179995.375000				
1381	1	indie triste	0.946000	0.83000
207400.000000				

1239	1	guaracha	0.009030	0.74500		
189818.000000						
2383	1	russian dance	0.005610	0.65300		
198095.000000						
	energy	instrumentalness	liveness	loudness	speechiness	\
237	0.630167	0.000020	0.081067	-6.627833	0.134833	
2755	0.777000	0.000035	0.086300	-4.246000	0.117000	
2533	0.562333	0.130339	0.075133	-7.719000	0.050733	
2778	0.719000	0.000004	0.326000	-4.923000	0.045500	
46	0.685000	0.000000	0.148000	-6.429000	0.062700	
536	0.787000	0.000000	0.323000	-4.592000	0.032000	
37	0.580187	0.000258	0.189950	-7.016687	0.195563	
1381	0.159000	0.000020	0.362000	-14.461000	0.038300	
1239	0.972000	0.465000	0.297000	-3.506000	0.077400	
2383	0.945000	0.915000	0.439000	-2.634000	0.096000	
	tempo	valence	popularity	key		
237	115.092500	0.588667	80.666667	2		
2755	102.020000	0.706000	80.000000	1		
2533	123.676333	0.834333	80.000000	1		
2778	120.062000	0.364000	80.000000	0		
46	99.954000	0.937000	78.500000	11		
536	142.018000	0.199000	78.500000	1		
37	83.250125	0.676625	77.312500	11		
1381	104.950000	0.189000	77.000000	1		
1239	128.031000	0.556000	77.000000	7		
2383	126.093000	0.326000	77.000000	5		
<i># Artist analysis</i>						
artist.sort_values('popularity', ascending=False).head(10)						
	mode	count	acousticness		artists	danceability
duration_ms	\					
20966	0	2	0.056300	Ritt Momney	0.399000	
210463.0						
14354	1	1	0.090700	Lele Pons	0.905000	
155825.0						
15070	0	2	0.310000	Los Legendarios	0.823000	
213314.0						
11764	1	2	0.819000	Jerry Di	0.854000	
197587.0						
7463	0	2	0.068600	Emilee	0.674000	
176547.0						
28263	1	2	0.424000	salem ilese	0.738000	
136839.0						
23687	0	1	0.068600	Surf Mesa	0.674000	
176547.0						
213	0	3	0.166633	A7S	0.742667	
168293.0						

14378	0	2	0.254000	Lenny Santos	0.752000	
206760.0						
16453	1	1	0.018500	Master KG	0.880000	
342613.0						
tempo	\	energy	instrumentalness	liveness	loudness	speechiness
20966	0.491		0.000890	0.110000	-10.778000	0.0538
91.066000						
14354	0.686		0.000000	0.266000	-3.152000	0.0664
103.013000						
15070	0.821		0.000004	0.143000	-3.402000	0.1660
99.999000						
11764	0.525		0.000000	0.146000	-4.426000	0.2140
97.054000						
7463	0.774		0.001880	0.393000	-7.567000	0.0892
112.050000						
28263	0.621		0.000007	0.692000	-7.313000	0.0486
113.968000						
23687	0.774		0.001880	0.393000	-7.567000	0.0892
112.050000						
213	0.726		0.000000	0.154667	-5.921333	0.1980
121.296333						
14378	0.716		0.000000	0.243000	-5.200000	0.0325
117.007000						
16453	0.483		0.000009	0.060700	-7.012000	0.0504
124.009000						
	valence	popularity	key			
20966	0.151000	93.0	6			
14354	0.963000	92.0	0			
15070	0.791000	90.0	8			
11764	0.630000	89.0	1			
7463	0.330000	88.0	11			
28263	0.715000	88.0	0			
23687	0.330000	88.0	11			
213	0.554667	87.0	8			
14378	0.553000	86.0	10			
16453	0.827000	86.0	1			

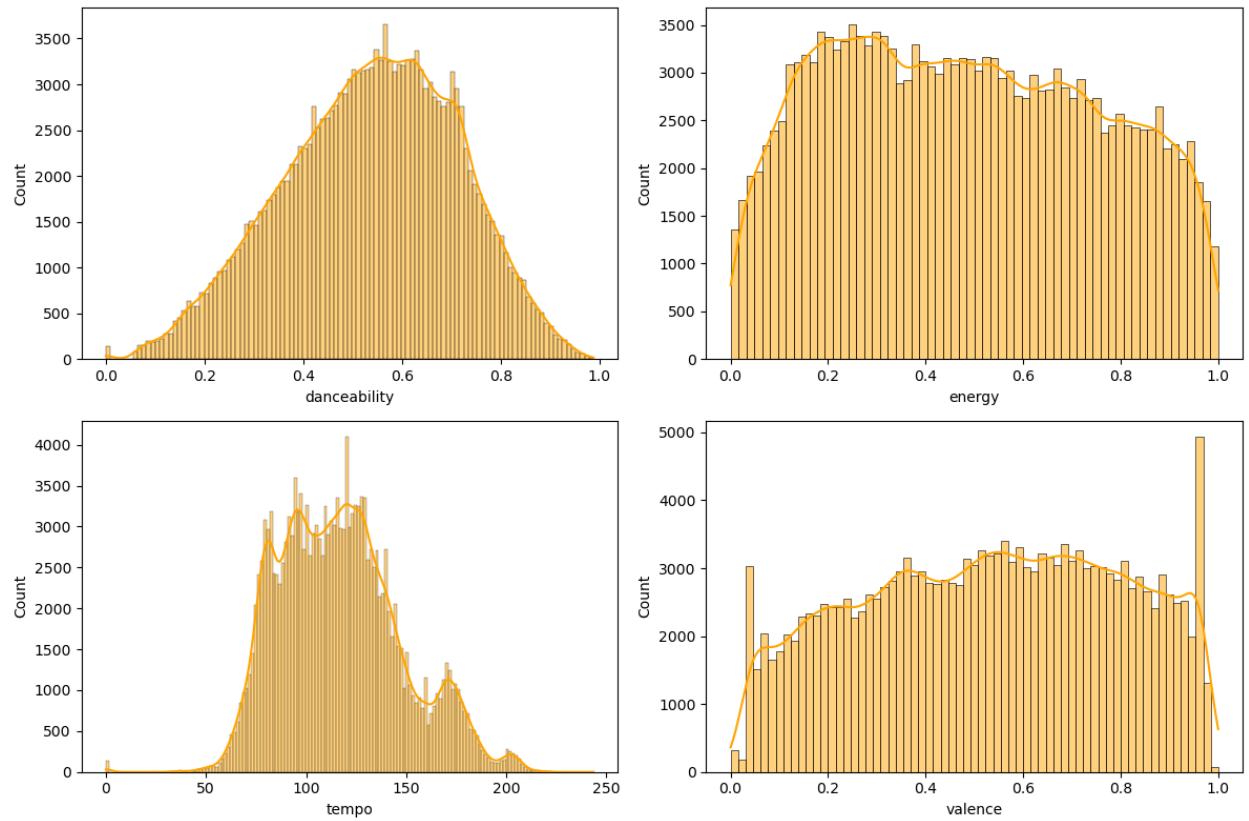
```
# Yearly trends
plt.figure(figsize=(12,3))
sns.barplot(data=year.tail(15),x='year',y='popularity',hue='popularity')
plt.xlabel("Year")
plt.ylabel("Average Popularity")
plt.show()
```



Data Analysis

Task 3.1: Distribution of Features (danceability, energy, tempo)

```
# Histograms with KDE
features = ['danceability', 'energy', 'tempo', 'valence']
plt.figure(figsize=(12,8))
for i, col in enumerate(features, 1):
    plt.subplot(2,2,i)
    sns.histplot(tracks[col], kde=True, color='orange')
plt.tight_layout()
plt.show()
```

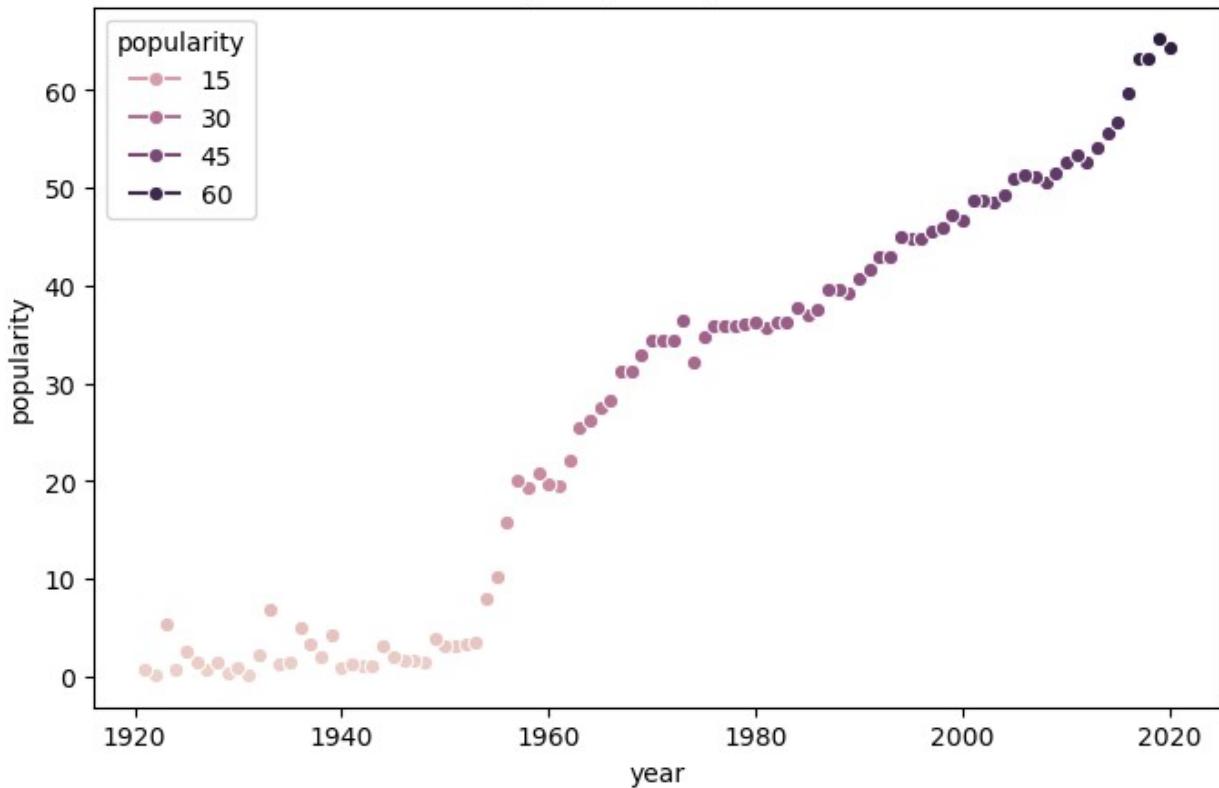


Task 3.2: Trends Over Time

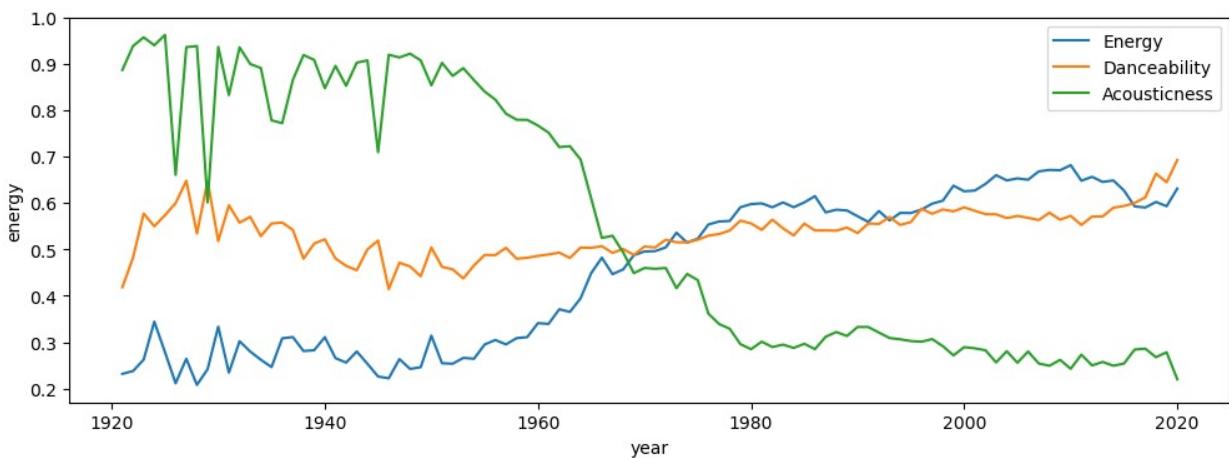
Using data_by_year.csv to observe how music changes historically.

```
# Popularity trend line
plt.figure(figsize=(8,5))
sns.lineplot(data=year,x='year',y='popularity',marker='o',hue='popularity')
plt.title("Average Popularity Over Time")
plt.show()
```

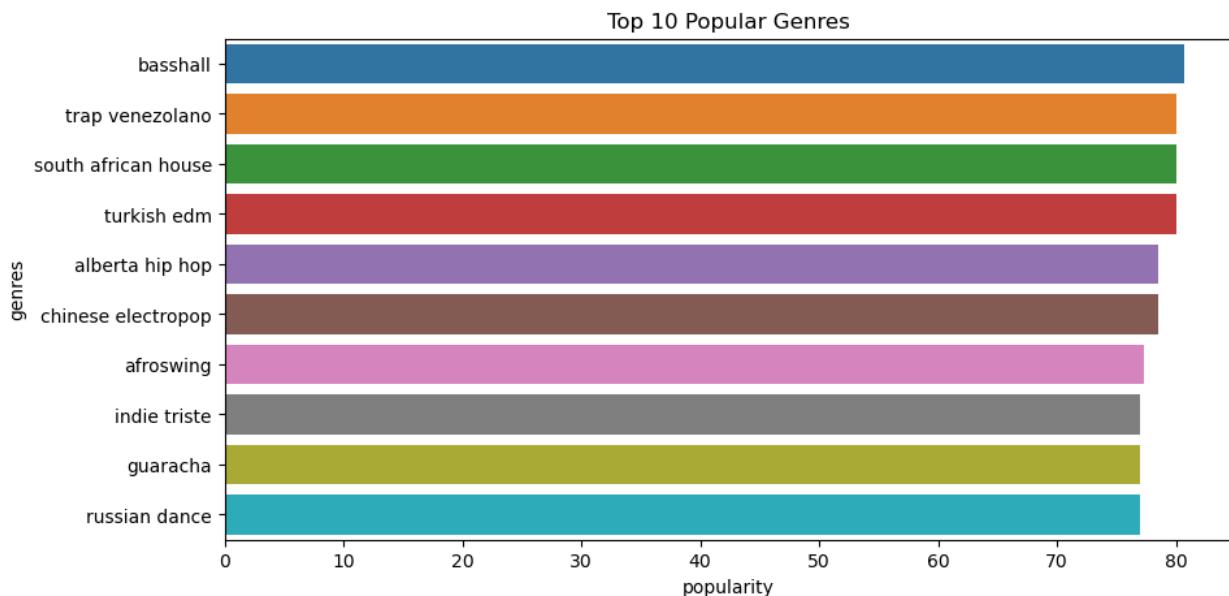
Average Popularity Over Time



```
# Feature evolution over time
plt.figure(figsize=(12,4))
sns.lineplot(data=year,x='year',y='energy',label='Energy')
sns.lineplot(data=year,x='year',y='danceability',label='Danceability')
sns.lineplot(data=year,x='year',y='acousticness',label='Acousticness')
plt.legend()
plt.show()
```

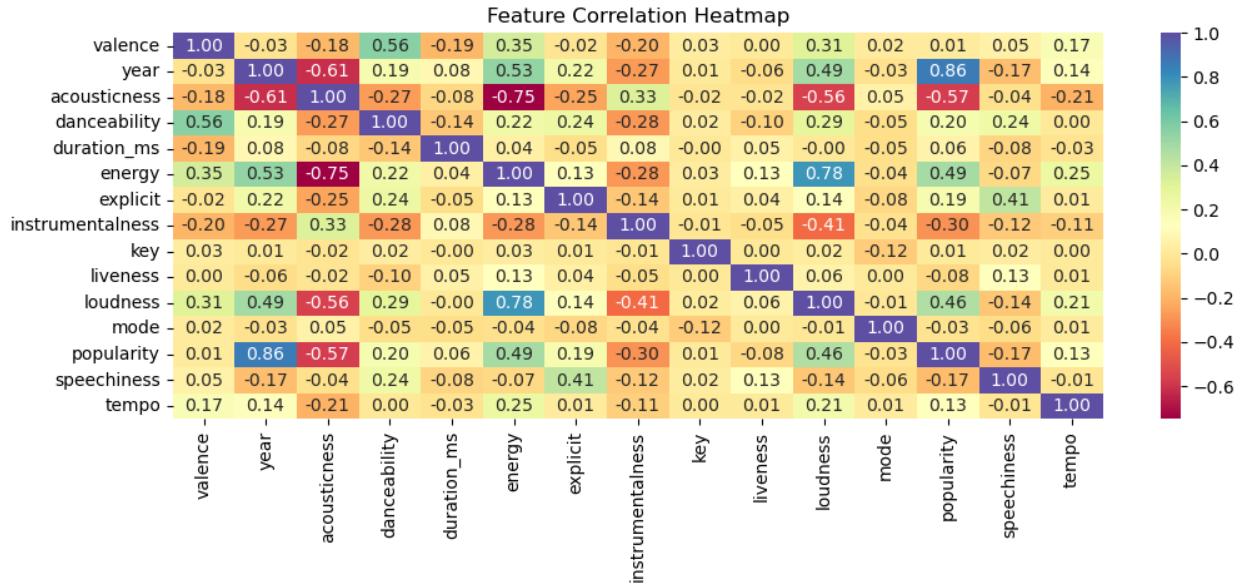


```
# Most popular genres
top_genres = genres.sort_values('popularity', ascending=False).head(10)
plt.figure(figsize=(10,5))
sns.barplot(data=top_genres, x='popularity', y='genres', hue='genres')
plt.title('Top 10 Popular Genres')
plt.show()
```

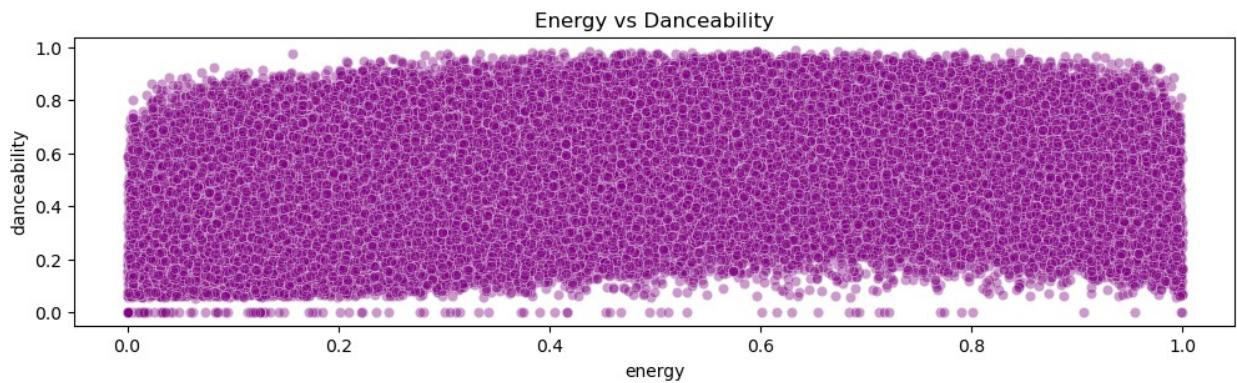


Task 3.3: Correlation Between Features

```
# Correlation heatmap
plt.figure(figsize=(12,4))
corr = tracks.corr(numeric_only=True)
sns.heatmap(corr, annot=True, cmap='Spectral', fmt='.2f')
plt.title('Feature Correlation Heatmap')
plt.show()
```



```
# Energy vs Danceability scatter
plt.figure(figsize=(12,3))
sns.scatterplot(data=tracks,x='energy',y='danceability',alpha=0.4,colo
r='purple')
plt.title('Energy vs Danceability')
plt.show()
```

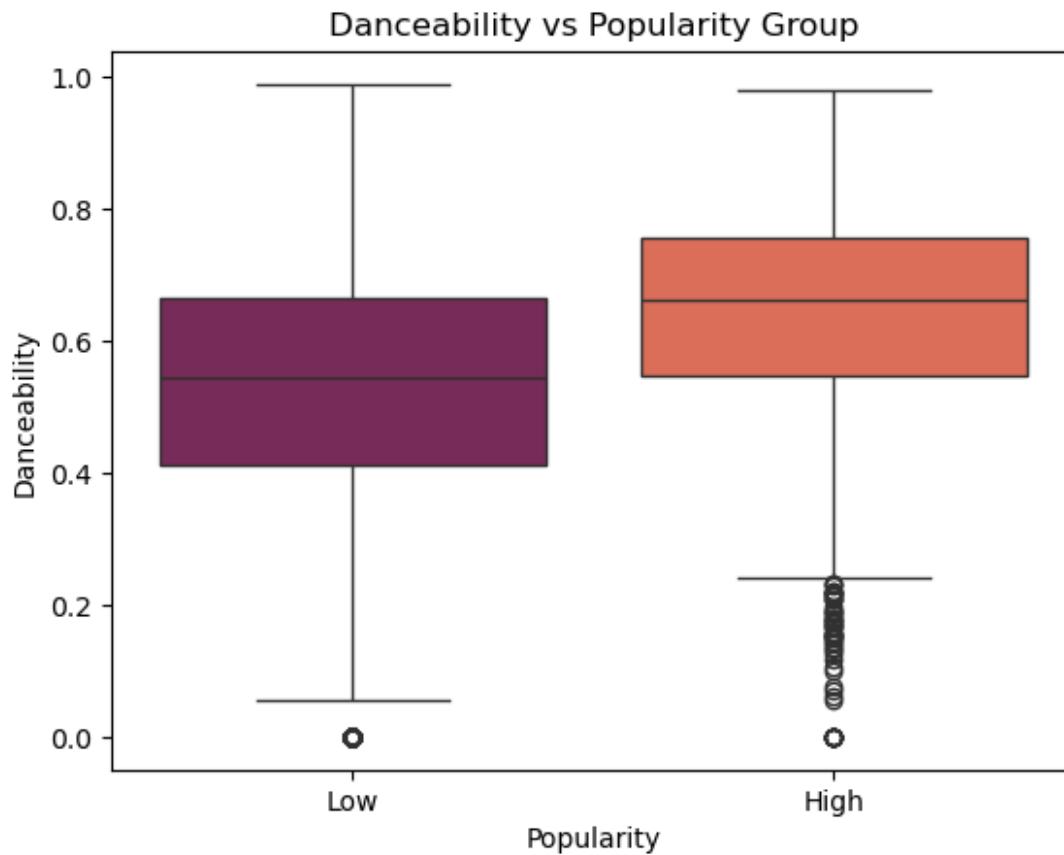


Task 3.4: User Preferences & Listening Habits

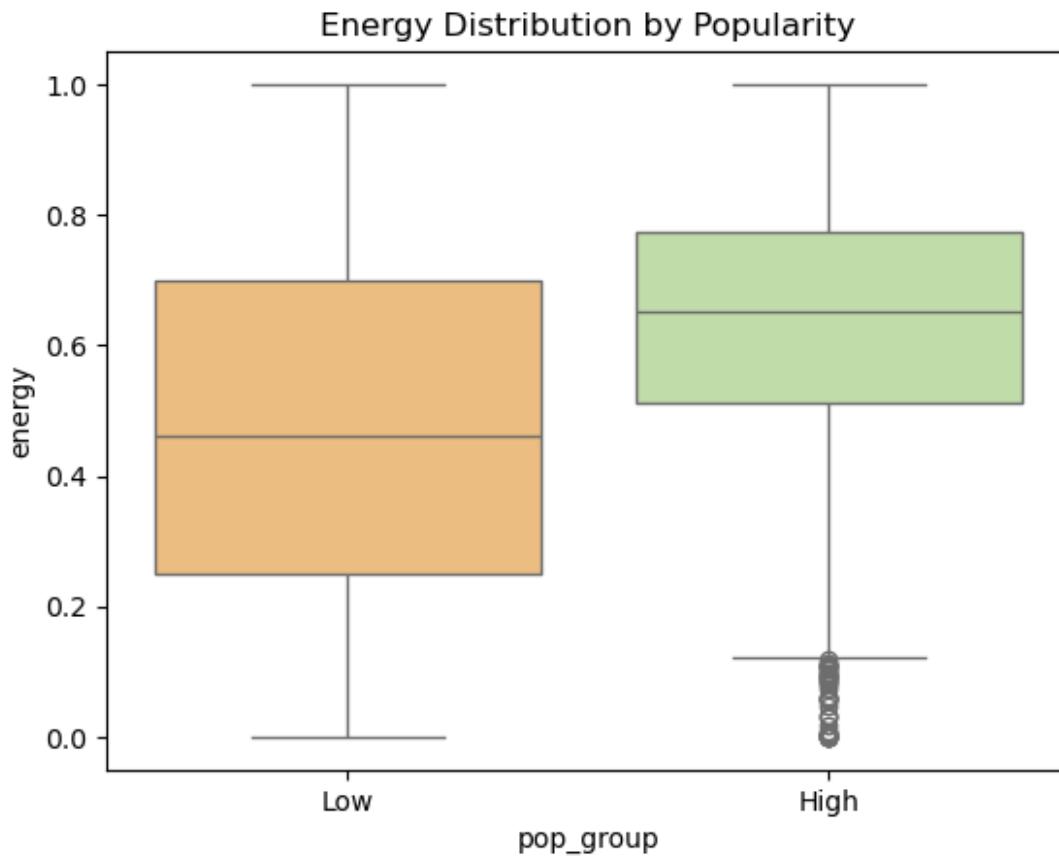
```
# Popular vs less popular songs (boxplot comparison)
tracks['pop_group'] = tracks['popularity'].apply(
    lambda x: 'High' if x >= 70 else 'Low')

sns.boxplot(data=tracks,x='pop_group',y='danceability',palette='rocket
')
plt.xlabel('Popularity')
plt.ylabel('Danceability')
```

```
plt.title('Danceability vs Popularity Group')
plt.show()
```



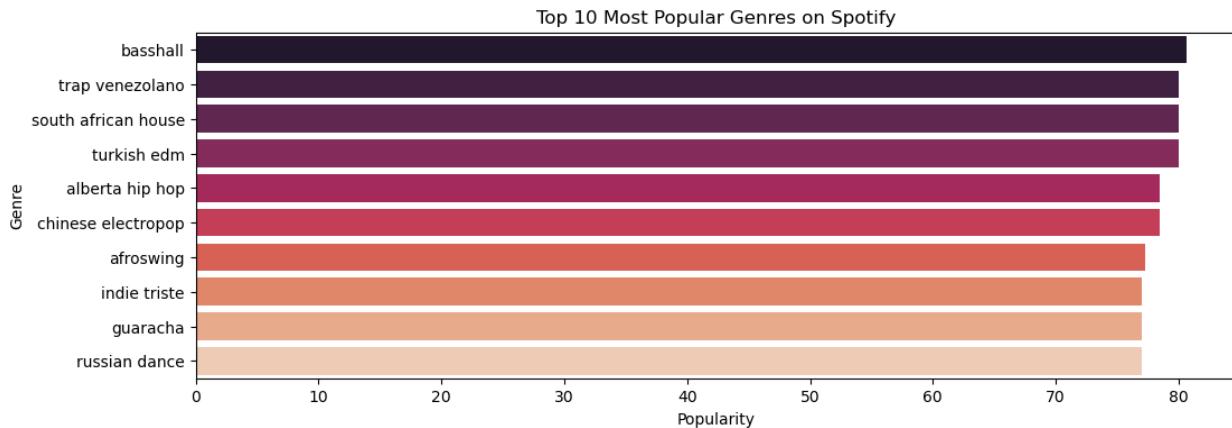
```
# Feature comparison for popularity
sns.boxplot(data=tracks,x='pop_group',y='energy',palette='Spectral')
plt.title("Energy Distribution by Popularity")
plt.show()
```



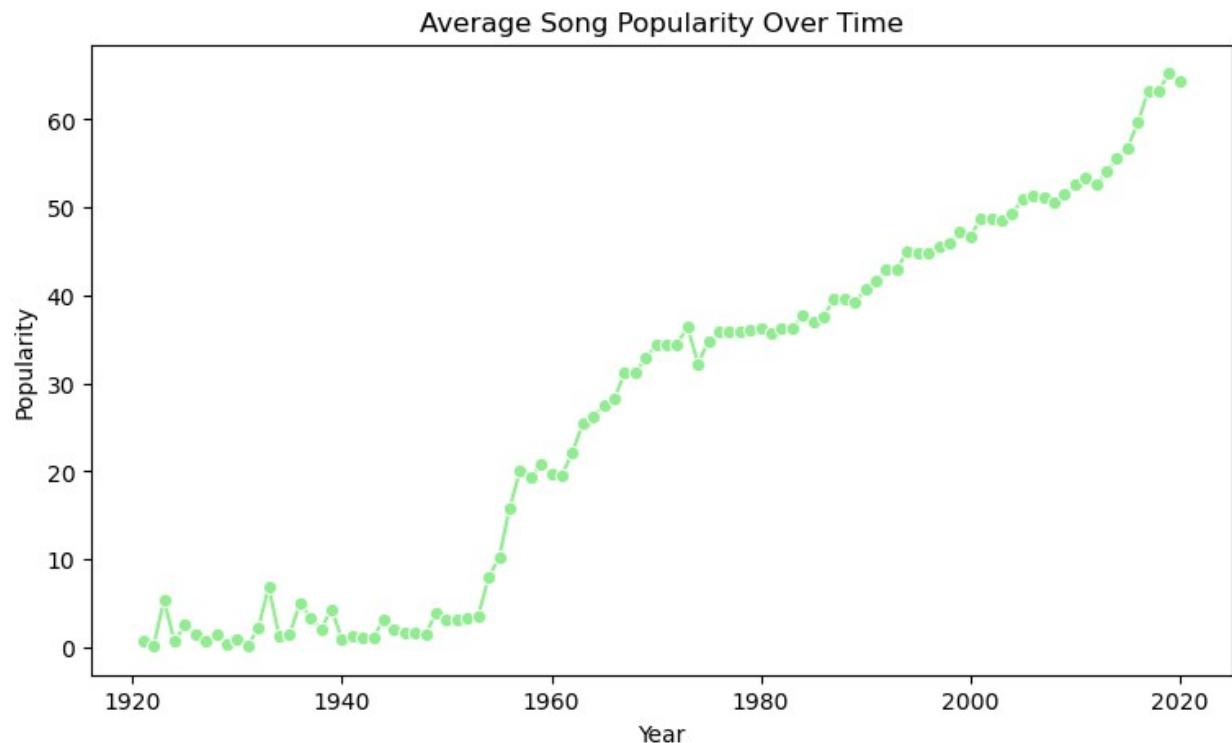
Visualization

Task 4.1: Basic Visualizations (Bar, Line, Scatter)

```
# Bar Chart – Top 10 Popular Genres
top_genres = genres.sort_values('popularity',
                                ascending=False).head(10)
plt.figure(figsize=(12,4))
sns.barplot(data=top_genres,x='popularity',y='genres',palette='rocket')
plt.title("Top 10 Most Popular Genres on Spotify")
plt.xlabel("Popularity")
plt.ylabel("Genre")
plt.show()
```

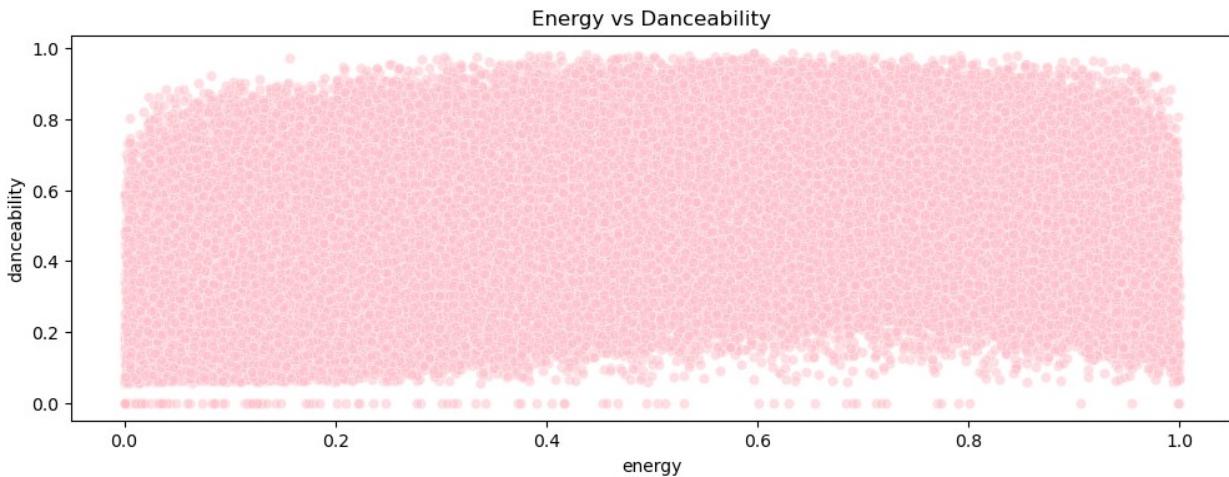


```
# Line Graph – Popularity Trend Over Time
plt.figure(figsize=(9,5))
sns.lineplot(data=year,x='year',y='popularity',marker='o',color='lightgreen')
plt.title("Average Song Popularity Over Time")
plt.xlabel("Year")
plt.ylabel("Popularity")
plt.show()
```



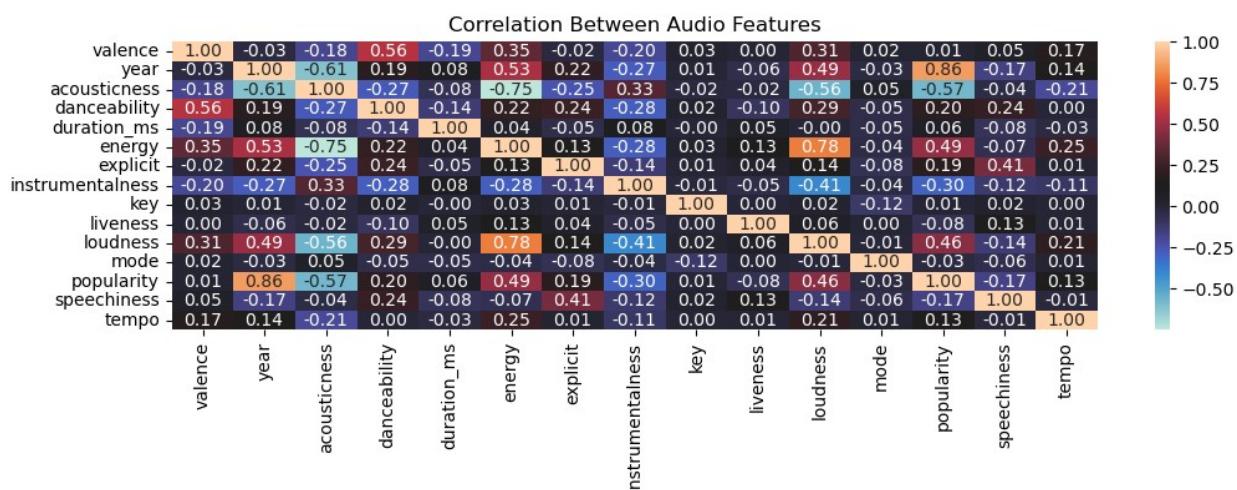
```
# Scatter Plot – Energy vs Danceability
plt.figure(figsize=(12,4))
sns.scatterplot(data=tracks,x='energy',y='danceability',alpha=0.5,colo
```

```
r='pink')
plt.title("Energy vs Danceability")
plt.show()
```



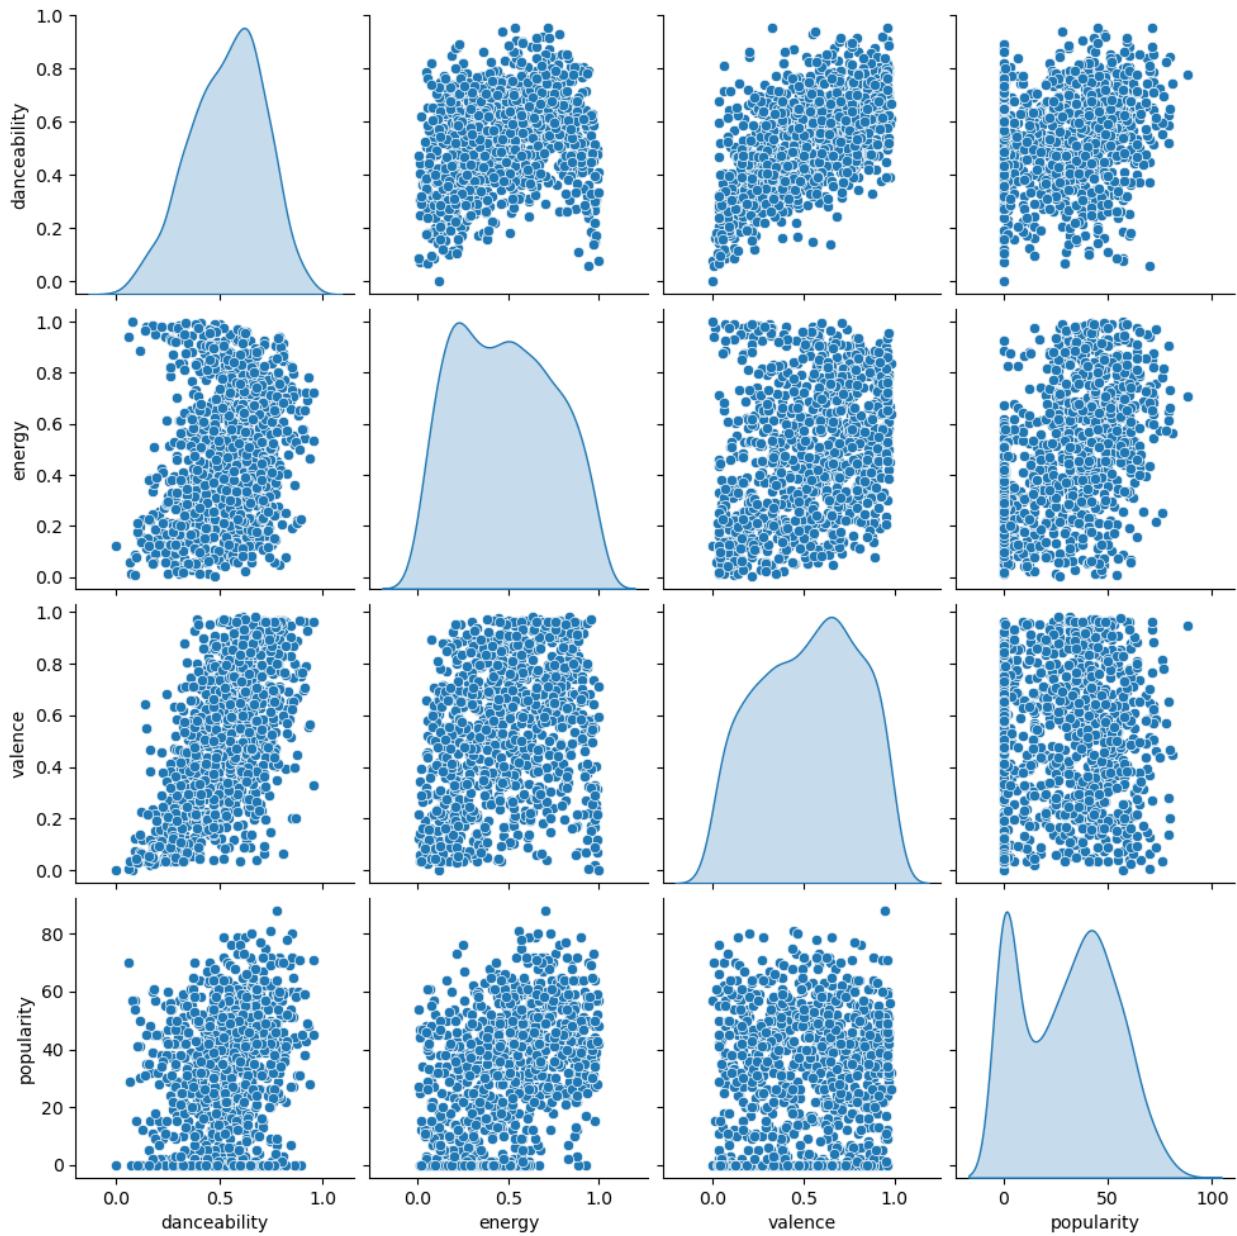
Task 4.2: Advanced Visualizations (Heatmaps, Pair Plots)

```
# Correlation Heatmap
plt.figure(figsize=(12,3))
corr = tracks.corr(numeric_only=True)
sns.heatmap(corr, annot=True, cmap='icefire', fmt='.2f')
plt.title('Correlation Between Audio Features')
plt.show()
```



```
# Pair Plot (Feature Relationships)
sample_tracks = tracks.sample(1000) # speed up plotting
```

```
sns.pairplot(sample_tracks[['danceability', 'energy', 'valence', 'popularity']], diag_kind='kde')
plt.show()
```



Modeling and Predictions

Task 5.1: Build Predictive Models to Forecast Song Popularity

```
# Prepare the data
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
# Select numeric features only
features = [
    'danceability', 'energy', 'loudness', 'speechiness',
    'acousticness', 'instrumentalness', 'liveness',
    'valence', 'tempo', 'duration_ms'
]

X = tracks[features]
y = tracks['popularity']
# Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42
)
# Scale features (important for regression)
scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_test = scaler.transform(X_test)
```

Task 5.2: Evaluate Different Models

```
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_absolute_error, mean_squared_error,
r2_score

lr = LinearRegression()
lr.fit(X_train, y_train)

y_pred_lr = lr.predict(X_test)

print("Linear Regression Results:")
print("MAE:", mean_absolute_error(y_test, y_pred_lr))
print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred_lr)))
print("R2:", r2_score(y_test, y_pred_lr))

Linear Regression Results:
MAE: 13.11202616522743
RMSE: 16.307952651171533
R2: 0.44422210532623396
```

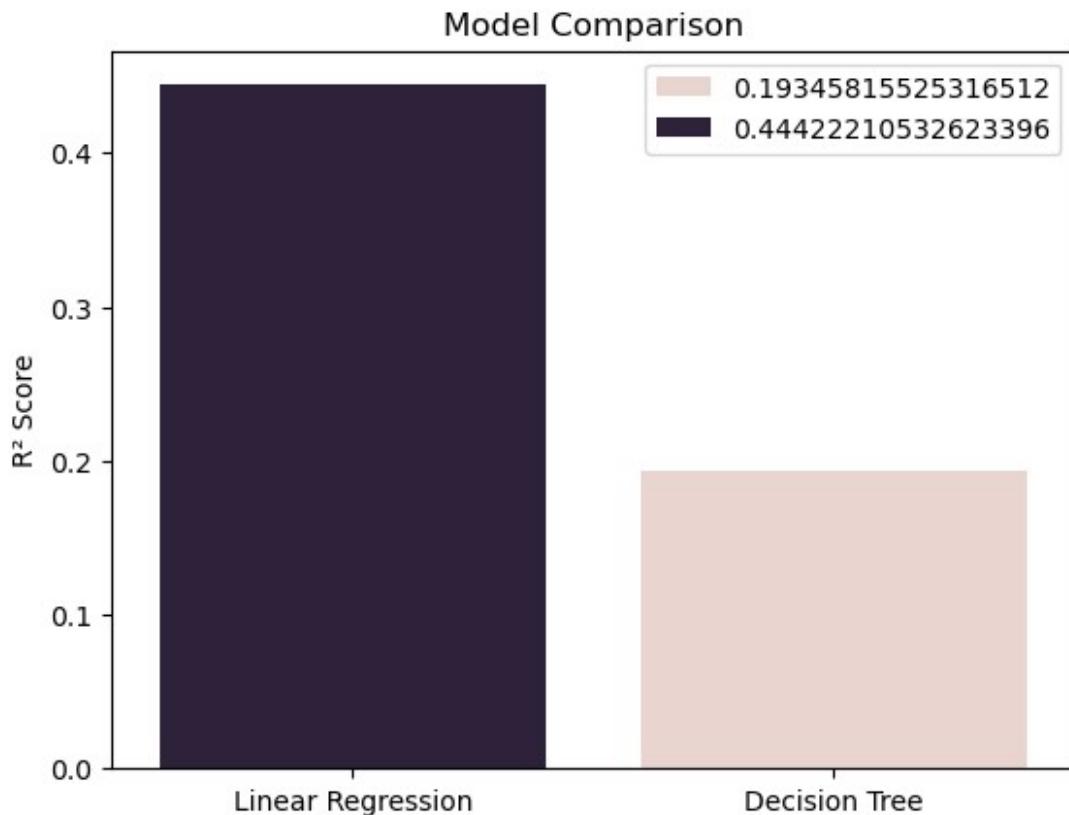
```
# Decision Tree
from sklearn.tree import DecisionTreeRegressor
dt = DecisionTreeRegressor(random_state=42)
dt.fit(X_train, y_train)

y_pred_dt = dt.predict(X_test)

print("Decision Tree Results:")
print("MAE:", mean_absolute_error(y_test, y_pred_dt))
print("RMSE:", np.sqrt(mean_squared_error(y_test, y_pred_dt)))
print("R2:", r2_score(y_test, y_pred_dt))

Decision Tree Results:
MAE: 14.255713769495962
RMSE: 19.645462578647304
R2: 0.19345815525316512

# Compare Model Performance (Visualization)
models = ['Linear Regression', 'Decision Tree']
r2_scores = [
    r2_score(y_test, y_pred_lr),
    r2_score(y_test, y_pred_dt)
]
sns.barplot(x=models,y=r2_scores,hue=r2_scores)
plt.ylabel("R2 Score")
plt.title("Model Comparison")
plt.show()
```



Task 5.3: Fine-Tune Models to Improve Accuracy

```
from sklearn.model_selection import GridSearchCV

params = {
    'max_depth': [5, 10, 20, None],
    'min_samples_split': [2, 5, 10]
}

grid = GridSearchCV(
    DecisionTreeRegressor(random_state=42),
    params,
    cv=5,
    scoring='r2'
)
grid.fit(X_train, y_train)

best_dt = grid.best_estimator_
y_pred_best = best_dt.predict(X_test)
```

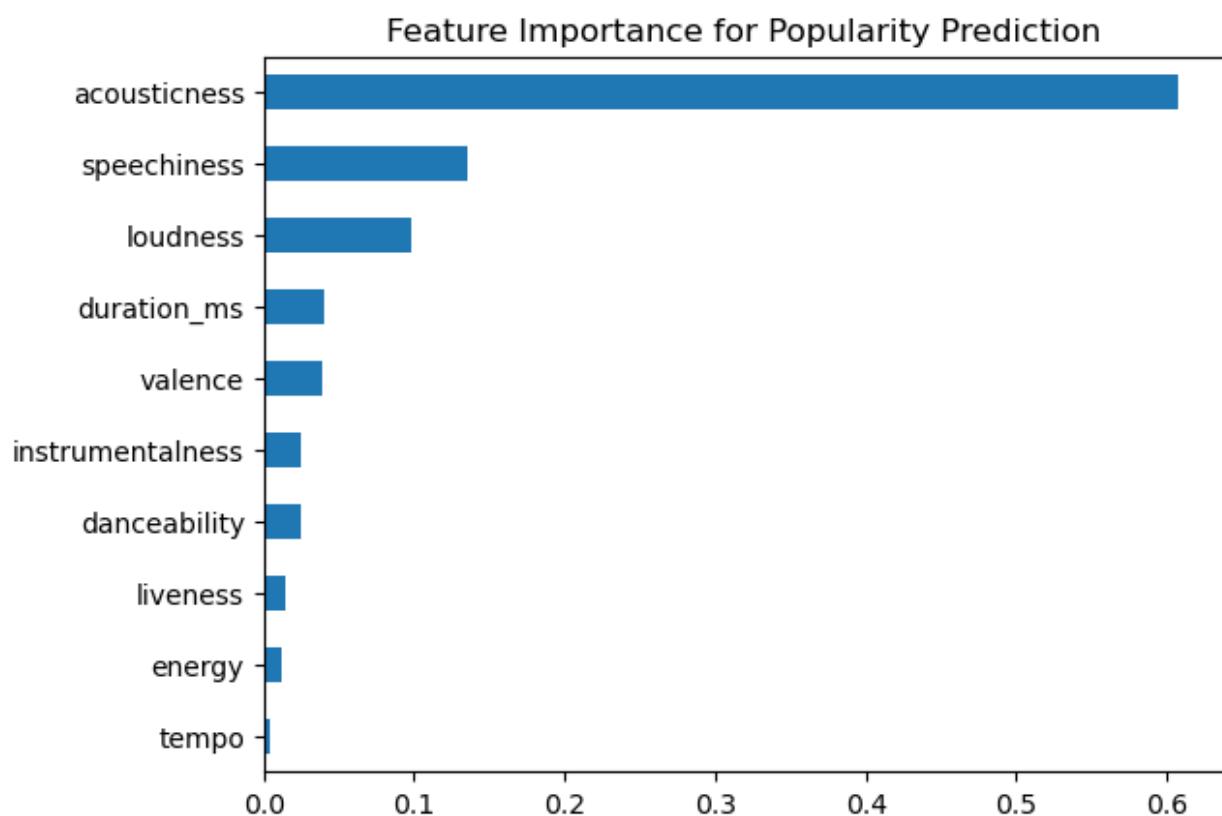
```

print("Best Parameters:", grid.best_params_)
print("Tuned R2:", r2_score(y_test, y_pred_best))

Best Parameters: {'max_depth': 10, 'min_samples_split': 10}
Tuned R2: 0.5415446259348499

# Feature Importance (Very useful insight)
importance = pd.Series(best_dt.feature_importances_, index=features)
importance.sort_values().plot(kind='barh')
plt.title("Feature Importance for Popularity Prediction")
plt.show()

```



Conclusion

Task 6.1: Key Findings

The analysis of Spotify's dataset revealed several important patterns in music trends, user preferences, and track characteristics. Popular tracks generally exhibit higher energy, danceability, and loudness, indicating that upbeat and engaging songs tend to attract more listeners. Certain genres consistently dominate streaming counts, suggesting strong audience loyalty toward specific music styles. Additionally, newer releases show higher engagement.

levels compared to older tracks, highlighting the impact of recency on streaming behavior. Correlation analysis also indicated relationships between audio features (such as energy and loudness) and popularity scores.

Task 6.2: Implications of the Results

These findings have practical implications for artists, producers, and music streaming platforms. Understanding which features contribute to track popularity can help artists tailor their music to audience preferences. Record labels can make data-driven decisions regarding genre investments and release strategies. For Spotify, insights into listener behavior can enhance recommendation systems, personalized playlists, and targeted promotions, ultimately improving user satisfaction and platform engagement.

Task 6.3: Future Research Directions

Future research could explore deeper user-level analytics, such as individual listening patterns and personalized recommendation accuracy. Time-series analysis may help identify evolving trends and seasonal listening habits. Incorporating sentiment analysis of lyrics or social media feedback could provide richer context for popularity prediction. Additionally, applying machine learning models to predict hit songs based on audio features would further enhance predictive capabilities.